

A Cluster-Local Softmax Merge Unit for RISC-V Vector Clusters

Anonymous Author(s)

Abstract—Blockwise attention updates a compact (m, l, O) state after every tile. A general-purpose core pays scalar control and normalization costs for every row, even when vector hardware handles the output update. This paper presents a cluster-local Softmax Merge Unit (SMU) for this recurrence. The SMU uses the existing memory-mapped control plane and TCDM scratchpad interface of a Spatz RISC-V vector cluster. It therefore leaves the ISA and the main vector pipeline unchanged. The datapath evaluates the maximum, two exponential scales, reciprocal-based weights, and the weighted output-vector update. Across 16 measured (N, D) coordinates, the full SMU outperforms an RTL-aligned scalar-plus-RVV baseline by 1.42–10.14 $\times$. Concurrent execution with a streaming core workload preserves the SMU execution time within 0.05% at (16, 64), while the core slows by 0.6%. The results show that a compact auxiliary engine can remove per-row merge overhead without requiring an ISA extension.

Index Terms—online softmax, attention, RISC-V vector processor, scratchpad, accelerator

I. INTRODUCTION

Tiled attention avoids materializing the full score matrix by updating online-softmax state after each tile [1], [2]. Each merge produces m_{new} , l_{new} , w_{old} , and w_{tile} before the regular $O[D]$ vector update. The same pattern repeats for every row and tile in the kernel. This decomposition separates the state-dependent scalar recurrence from uniform elementwise work. The recurrence has row-level dependencies, whereas the vector update exposes element-level parallelism, so the two phases have different dependency/parallelism structures.

Spatz provides compact RVV execution and shared TCDM, which suits the regular elementwise update [3]. We use B2R as the strong software baseline: RVV executes the $O[D]$ update while software executes the scalar recurrence. Maximum selection, exponential rescaling, length accumulation, and weight generation remain serial within each row and tile. The Full-Offload Ablation lowers scalar cost but replaces the efficient existing RVV datapath with a specialized vector datapath. This contrast makes the specialization boundary the design question.

Our proposed design is exactly Scalar SMU + existing RVV. The Scalar SMU executes the state-dependent recurrence and writes m_{new} , l_{new} , w_{old} , and w_{tile} to TCDM. After completion, software loads the two weights from TCDM and invokes existing RVV for the regular $O[D]$ update. This implementation leaves the existing ISA and RVV datapath unchanged (Figure 1).

Across three model-derived configurations, Proposed achieves a 4.81 $\times$ geometric-mean measured-cycle speedup over B2R, with speedup defined as B2R cycles divided by

Proposed cycles. Relative to B2R, Proposed reduces the fitted per-row scalar term C_s by 17.66 $\times$ while increasing the per-element vector term C_v by 1.88%; the Full-Offload Ablation lowers C_s further, but its C_v is 3.78 $\times$ Proposed and its standalone mapped SMU area is 48.8% larger. We make three contributions:

- **Lightweight Scalar SMU.** We design a cluster-local accelerator for the state-dependent online-softmax recurrence, covering maximum selection, exponential rescaling, length accumulation, reciprocal normalization, and weight generation.
- **Selective Offloading.** We assign the scalar recurrence to the SMU and retain the regular $O[D]$ update on the existing RVV datapath through a post-completion, TCDM-mediated weight handoff.
- **Quantitative Validation.** We measure RTL/Verilator cycles across scaling and model-derived configurations and report standalone SMU mapped cells and Nangate45 Liberty cell-area units from Yosys/ABC.

II. SELECTIVE-OFFLOAD ARCHITECTURE

Proposed is a Scalar SMU paired with existing RVV: the SMU handles the state-dependent scalar recurrence, while the core retains the regular $O[D]$ vector update. This selective boundary moves normalization control across the accelerator interface and leaves the established vector datapath and software sequencing in place. Shared TCDM carries the handoff state, so the architecture composes an RTL-defined scalar engine with an unchanged RVV update path. The scalar unit owns quantities whose next value depends on the row's selected maximum and normalization state; RVV owns operations that repeat over independent vector elements. The handoff therefore preserves software-visible memory ordering while narrowing hardware specialization to the recurrence. Figure 1 shows this selective-offload boundary and its TCDM-mediated handoff.

A. Specialization Boundary

Online merge separates state selection and normalization from the regular elementwise vector update. For one row, the




Fig. 1. Proposed architecture. Mode 1 directs the Scalar SMU to read scalar state and write m_{new} , l_{new} , w_{old} , and w_{tile} to shared TCDM. Software waits for completion, then invokes existing RVV for the $O[D]$ update; the diagram has no direct SMU-to-RVV link.

decomposition is

$$\begin{aligned}
 m_{new} &= \max(m_{old}, m_{tile}), & (1) \\
 s_{old} &= l_{old} \exp(m_{old} - m_{new}), & (2) \\
 s_{tile} &= l_{tile} \exp(m_{tile} - m_{new}), & (3) \\
 l_{new} &= s_{old} + s_{tile}, & (4) \\
 w_{old} &= s_{old} / l_{new}, \quad w_{tile} = s_{tile} / l_{new}, & (5) \\
 O_{new}[j] &= w_{old} O_{old}[j] + w_{tile} O_{tile}[j], \quad 0 \leq j < D. & (6)
 \end{aligned}$$

The first five lines execute in Scalar SMU, while the final line executes in existing RVV. EXP and reciprocal blocks implement RTL approximations, so these equations describe the decomposition rather than an IEEE-exact arithmetic claim. The index constraint makes the handoff explicit: the SMU produces per-row weights and RVV applies them across D elements. These lines reduce each row to two weights before any vector element moves. The decomposition also separates row-level dependencies from element-level accumulation, which defines the architecture's scheduling boundary.

Scalar recurrence belongs on a state-dependent datapath because both exponential arguments depend on the selected maximum and the reciprocal depends on the newly accumulated length. A regular vector loop excels at independent $O[j]$ operations after weights are known, but it does not remove the serial decisions that produce those weights. The boundary therefore assigns control-sensitive work to Scalar SMU and keeps uniform multiply-add work on existing RVV. This division avoids duplicating vector functional units while matching each phase to its natural execution granularity. Both decisions serialize within a row, whereas elements within

one O row share the resulting weights and expose regular parallelism. Keeping this distinction explicit lets the core reuse its tuned VLSU and FMA sequence.

B. Scalar Datapath and Schedule

The Scalar SMU implements the recurrence with a compact scalar datapath. A comparator selects m_{new} and forms two deltas; two EXP approximation/LUT blocks transform those deltas into scaled factors. Fixed-point multipliers combine the factors with l_{old} and l_{tile} , an accumulator forms l_{new} , and a reciprocal approximation normalizes both scaled lengths. Weight generation then emits w_{old} and w_{tile} for TCDM writeback. The RTL places these operations in the scalar states. In mode 1, control exits after scalar writeback and never enters the engine's UPDATE_VECTOR state; Proposed performs the vector update on existing RVV. All intermediate values remain on the scalar path until writeback, so the datapath follows the recurrence's data dependencies. The resulting interface exports weights as ordinary TCDM state for the next phase.

RTL observer counts establish a fixed per-row scalar schedule: LOAD_SCALAR consumes 13 cycles, COMPUTE_SCALAR 1, COMPUTE_WEIGHT 1, and STORE_SCALAR 9, for 24 SMU-busy cycles per row. LOAD and STORE account for 22 of 24 cycles as TCDM communication states; the two compute states occupy the remaining two cycles. The 24-cycle count describes the scalar-only FSM schedule. It is neither the fitted $C_s = 90.28$ nor software-visible or end-to-end latency, which also includes command sequencing and completion polling. The observer records only the four scalar states because mode 1 never enters UPDATE_VECTOR. The communication count

reflects four reads plus four writes at state-machine handshake granularity, while the compute states represent combinational evaluations.

C. TCDM-Mediated RVV Handoff

Mode 1 makes the handoff state explicit in TCDM. For each row, the Scalar SMU reads four FP32 scalar inputs, m_{old} , l_{old} , m_{tile} , and l_{tile} , from their TCDM arrays. It computes the new scalar state and writes m_{new} , l_{new} , w_{old} , and w_{tile} back to TCDM for all rows. The output weights are software-visible arrays rather than private accelerator state, which gives the core a stable handoff point. Because the weight arrays share TCDM with the state, the SMU does not retain a private weight interface for the core. The same row indexing keeps scalar outputs aligned with subsequent vector rows.

Software sequences the two engines through status and memory. It programs the MMIO command, starts mode 1, polls the status until done, and then reads the two weight arrays from TCDM. After completion, software invokes existing RVV to load $O_{old}[D]$ and $O_{tile}[D]$, apply the weights, and store $O_{new}[D]$. The ordering gives RVV a completed, software-addressable weight state; no direct SMU-to-RVV path or weight-register bypass exists. Software remains the sequencing agent: it observes completion before exposing weights to RVV, so the core never consumes partially written rows. The RVV routine retains its existing load, multiply, FMA, and store sequence.

MMIO supplies command and status signals without changing the ISA. The cluster peripheral presents the software-selected operation to the Scalar SMU, while the engine contributes an independent TCDM requester to the shared interconnect. Existing core ports continue to serve RVV loads and stores, so the handoff traverses shared TCDM rather than a private engine link. This structure changes neither instruction decoding nor the RVV datapath. The interconnect arbitrates both requester classes through existing TCDM machinery, while MMIO remains the control-plane boundary. This composition adds no decoder instruction or custom vector operand.

This boundary fixes the architecture evaluated next: Scalar SMU removes state-dependent scalar recurrence work, and existing RVV retains the regular vector update through a software-mediated TCDM handoff. The Evaluation section measures the resulting cycle behavior and then separates scaling, model-derived configurations, and standalone hardware evidence. This framing lets the next section attribute cycle changes to the scalar recurrence while retaining a fixed vector implementation in the primary comparison. It also keeps standalone evidence separate from measured cycles.

III. EVALUATION

Selective scalar offloading retains the low vector cost of existing RVV and yields a $4.81 \times$ geometric-mean speedup in measured cycles over B2R. The Proposed design combines a Scalar SMU with existing RVV, while Full serves only as the Full-Offload Ablation. B2R supplies the head-to-head reference because it already performs the vector update with

existing RVV. This reference makes scalar recurrence offload the measured question and keeps Full’s role diagnostic. The primary result therefore concerns selective offload with the vector path held fixed in the comparison.

A. Experimental Setup

We evaluate FP32 state on a Spatz cluster with the (m, l, O) arrays resident in TCDM. The simulator runs RTL through Verilator, and every reported cycle count comes from the measured RTL/Verilator run. Each implementation uses the same generated state and command shape, so cycle differences reflect the execution path. The state includes old and tile scalar values together with output vectors, and the benchmark places source and destination arrays in TCDM before launch. We use one cluster configuration and one launch protocol across roles. RTL/Verilator cycles include the software-visible command and completion work, so each role exposes the path that its implementation actually executes.

We compare four roles. B1 runs the scalar software context. B2R runs the scalar recurrence in software and uses existing RVV for the vector update, making it the primary strong baseline. Proposed combines Scalar SMU with existing RVV. Full runs the Full-Offload Ablation, which moves the vector update into the SMU as well. We retain only outputs that pass the correctness gate; incomplete or failing cases do not enter a primary result. B1 supplies context for total scalar work but supplies no primary speedup. B2R keeps the vector loop on the existing RVV path, so its cycle count establishes the cost after vectorization. Proposed moves only the scalar recurrence across the SMU boundary, whereas Full moves both recurrence and vector update. This role assignment keeps each comparison tied to one architectural change.

We use Yosys/ABC with the Nangate45 library to collect standalone hardware evidence. An area-oriented `abc -fast` mapping reports formal cell counts and mapped area. A separate timing-driven mapping reports critical delay and an estimated standalone Fmax. The two flows answer different questions, so we keep their cell and area outputs separate from timing outputs. The available records provide no cluster PPA or power measurement. The area flow supports a standalone size comparison, while the timing flow supports a standalone timing estimate under its own constraints. We never combine a timing-flow cell count with the area-flow result. The evidence therefore supports the SMU-local comparison and leaves cluster-level PPA and power outside the reported scope.

B. Progressive Baseline

B1 to B2R isolates the contribution of RVV vectorization. B1 remains in Table I as cycle context, while B2R supplies the primary strong baseline for every speedup. This comparison identifies the cost that existing RVV removes before any SMU offload. The B1 column lets readers see how the scalar context grows with each model-derived shape, but it carries no speedup column. B2R preserves the same state representation and command setup while using the existing vector instructions, which makes the next comparison isolate the recurrence itself.

Every primary speedup divides B2R cycles by Proposed cycles. The ratio attributes improvement to scalar recurrence offload after RVV handles the vector update. Full appears only to quantify the cost and benefit of moving that vector update as an ablation. A common denominator keeps the interpretation stable across shapes and prevents a scalar-context comparison from becoming a headline result. The reported ratios thus measure the incremental contribution of Scalar SMU under the same RVV vector path.

C. Scaling and Cost Decomposition

To explain the cycle trend, we fit $C = C_0 + C_s N + C_v ND$ to each design. Here N denotes rows, D denotes vector dimension, C_0 captures fixed work, C_s captures per-row scalar recurrence work, and C_v captures per-element vector work. The fitted coefficients summarize RTL/Verilator cycle trends; they do not equal FSM busy-cycle counts. Figure 2 presents the fit and uses a logarithmic scale for C_s . The decomposition treats N and D as independent shape variables, so each coefficient summarizes a different growth direction. C_0 absorbs launch and loop-independent work, whereas C_s and C_v describe fitted contributions across the measured coordinates. None of these coefficients is a direct FSM state counter.

B2R fits $C_s/C_v = 1594.23/2.09$, whereas Proposed fits $90.28/2.13$. Proposed therefore cuts the scalar coefficient by $17.66\times$ while increasing the vector coefficient by only 1.88% . The near-constant C_v shows that Scalar SMU offload removes recurrence work after RVV vectorization without changing the existing vector path. The fitted scalar term accounts for the large difference between the two paths, while the fitted vector term stays within the same scale. This separation links the measured speedup to the specific recurrence work that Proposed moves off the core.

Full fits $C_s/C_v = 19.32/8.03$, and its C_v is $3.78\times$ Proposed's. Full therefore reduces the scalar coefficient further but pays a larger vector coefficient when it absorbs the vector update. The decomposition points to the selective tradeoff: low C_s with low C_v gives Proposed its measured advantage. Full quantifies the cost of extending scalar offload into the vector path, so the ablation separates the two sources of cycle change. Proposed keeps the lower vector coefficient while retaining the scalar reduction that matters for the primary baseline.

D. Model-Derived Configurations

We next evaluate one attention-position shape per model. For each shape, we generate a deterministic merge state with FP32 arrays in TCDM and measure the four roles under the same RTL/Verilator flow. The selected shapes are BERT $(N, D) = (12, 64)$, Mistral $(32, 128)$, and Qwen14B $(40, 128)$. Table I reports the resulting cycles, with B1 included as context. Each row represents one attention-position merge state, so the table exercises the recurrence at a model-derived (N, D) shape. Deterministic state generation keeps the inputs aligned across roles and makes the measured cycle columns directly comparable. The table reports merge cycles only and uses no end-to-end inference schedule.

The Proposed/B2R speedups are $5.04\times$ for BERT, $4.62\times$ for Mistral, and $4.77\times$ for Qwen14B. Their geometric mean is $4.81\times$, matching the section result. These ratios keep the vector update in both compared paths and isolate scalar recurrence offload across model-derived dimensions. The three ratios remain close because each workload uses the same role definitions and the same vector path in Proposed and B2R. The geomean therefore summarizes a consistent selective-offload effect across the tested model dimensions rather than a single favorable shape.

Qwen72B's selected shape exceeds local TCDM capacity, so the benchmark has no cycle result for that shape and this boundary closes the measured feasible scope.

E. Hardware Cost and Full-Offload Ablation

Table II separates area evidence from timing evidence. The area flow uses formal cell count and mapped area from abc -fast; the timing flow uses a separate timing-driven mapping for critical delay and standalone Fmax. We keep timing-flow cell and area outputs out of the area comparison so the two mappings answer their intended questions. The table consequently reports formal cells and area from one flow, then delay and estimated Fmax from the other. This separation preserves the meaning of each column and prevents a timing result from changing the area comparison.

Proposed maps to 73,505 cells and 77,103.3 Liberty units with a 12.34ns critical delay and an estimated standalone Fmax of 81.02 MHz. Full maps to 107,372 cells and 114,713.3 Liberty units with a 13.20ns critical delay and an estimated standalone Fmax of 75.74 MHz. Full increases standalone mapped area by 48.8% over Proposed. The Proposed row therefore combines the smaller formal mapped design with the shorter timing-flow delay in this standalone evidence. Full exposes the additional mapped logic required when the vector update joins the offload path.

Figure 3 places B2R at $1.00\times$ speedup with zero SMU increment. Proposed reaches $4.81\times$ at 77.103k mapped area, and Full reaches $2.00\times$ at 114.713k. B2R's zero SMU increment denotes the software baseline and does not denote zero cluster area. The x-coordinate tracks standalone SMU mapped area, while the y-coordinate tracks measured cycles relative to B2R. The Full point shows how extra vector offload changes this standalone tradeoff, and workload crosses show variation across all three model-derived shapes in the measured comparison.

Using a common 81.0186 MHz assumption yields a $2.40\times$ Proposed/Full derived throughput ratio for the two standalone designs. The same frequency gives a $3.57\times$ incremental standalone-SMU area-efficiency ratio. These A1/A2-only derived comparisons do not measure system throughput or cluster efficiency. Cluster area, cluster timing and Fmax, physical area, signoff, power, and energy remain unavailable. The selective tradeoff is clear: Scalar SMU plus existing RVV preserves low vector cost while Full spends area to offload the vector update. The frequency assumption lets us compare the two standalone designs on one common scale, but it cannot turn




Fig. 2. Fitted C_s and C_v from RTL/Verilator cycles for B2R, Proposed, and Full. The C_s axis uses a log scale. The fitted terms describe cycle trends and exclude wall-time and FSM-busy-cycle interpretations.

TABLE I
MODEL-DERIVED ONE-POSITION SHAPES AND MEASURED RTL/VERILATOR CYCLES. VALUES DESCRIBE THE MERGE MICROBENCHMARK, NOT END-TO-END INFERENCE. B1 PROVIDES CONTEXT AND HAS NO SPEEDUP COLUMN.

Workload	(N, D)	B1 context	B2R	Proposed	Full ablation	B2R/Proposed
BERT	(12, 64)	271,718	20,785	4,128	7,704	$5.04\times$
Mistral	(32, 128)	1,401,702	59,500	12,866	34,728	$4.62\times$
Qwen14B	(40, 128)	1,780,547	75,027	15,733	43,206	$4.77\times$
Geomean	—	—	—	—	—	$4.81\times$

TABLE II
STANDALONE MAPPED HARDWARE EVIDENCE. THE AREA AND TIMING COLUMNS USE SEPARATE MAPPINGS; TIMING-FLOW CELLS AND AREA DO NOT ENTER THE AREA FLOW. FMAX IS A PRE-LAYOUT STANDALONE ESTIMATE.

Design	Mapped cells (area flow)	Mapped area (Liberty units)	Critical delay (timing flow)	Est. standalone Fmax
Proposed	73,505	77,103.3	12.34 ns	81.02 MHz
Full-Offload Ablation	107,372	114,713.3	13.20 ns	75.74 MHz

mapped evidence into a system measurement. The conclusion stays within A1/A2 evidence: selective scalar offload retains vector cost discipline while Full trades additional standalone area for vector-path offload in this ablation.

IV. DISCUSSION

A. The design targets a specific granularity

The SMU receives a complete merge command, so each launch carries a fixed configuration and completion cost. The measurements show that this choice works best when several rows amortize that cost. At (1, 1), only 29 SMU busy cycles execute inside a 1,124-cycle end-to-end operation. At (16, 64), the engine executes 8,534 busy cycles inside 9,556 end-to-end cycles. This change explains why the engine remains useful

for matrix shapes that contain many short scalar recurrences, even when each recurrence has a compact vector.

The command boundary also defines a clear deployment path. A tiled attention kernel can retain its QK and PV computation on an existing vector or matrix engine, place the partial (m, l, O) state in local scratchpad memory, and submit merge commands at tile completion. The current prototype serializes one command, but independent rows and heads expose natural opportunities for multiple engines. A future front end can batch commands or use a queue to reduce the small-workload overhead that the current MMIO protocol reveals.

B. The evaluation distinguishes evidence layers

The paper uses three measurement layers because each answers a separate question. RTL end-to-end cycles establish the




Fig. 3. Standalone SMU mapped area on the x-axis and frequency-independent measured-cycle speedup over B2R on the y-axis. B2R’s zero SMU increment does not denote zero cluster area. Circles mark geomeans and crosses mark workloads; the plot makes no cluster-area claim.

software-versus-engine result. FSM counters establish where the engine spends active cycles. The concurrency experiment establishes whether a co-running stream alters those component times. The combination gives a stronger explanation than a speedup number alone: the scalar work leaves the core, the vector loop becomes dominant, and the shared scratchpad introduces measurable yet small interference for the tested stream.

The numerical results also have two layers. The host model tests the selected LUT interpolation and reciprocal construction against the complete recurrence. The RTL benchmark tests the implemented fixed-point data path against its RTL-aligned software reference. Reporting both values prevents the lower model error from becoming an end-to-end RTL claim. The explicit validation paths serve the same purpose. They show that the engine rejects malformed commands instead of silently producing an output.

C. Proxy data guides follow-up work

The generic logic and toggle measurements guide architectural choices but do not replace physical implementation. The resource decomposition assigns most generic cells to the vector arithmetic and data path, while the exponential and reciprocal blocks occupy smaller independent scopes. The toggle traces show a lower count for B3 at the largest

sampled coordinate because B3 completes the operation in fewer target cycles. These observations motivate vector-path optimization and physical synthesis. They do not establish an area, frequency, power, or energy advantage.

V. RELATED WORK

Dedicated Softmax and attention accelerators offload broader Softmax or attention datapaths. Softermax combines unnormalized and normalization hardware with online normalization for Transformer softmax [4]. Hyft provides a reconfigurable full-vector Softmax pipeline for training and inference [5]. SoftEx provides the closest shared-TCDM/HWPE precedent, but it still offloads the full Softmax path [6]. Alexandridis et al. add custom exponent and multiplication hardware to FlashAttention, including recurrence-related computation and the output-vector update [7]. These designs accelerate a complete Softmax or attention stage; they do not isolate the scalar merge recurrence from the regular vector update.

Programmable RVV systems expose a different point in the design space. Spatz provides the RVV/TCDM substrate used by our cluster [3]. Spatzformer reconfigures a dual-core RVV cluster between split and merge modes, using generic mixed scalar-vector resource reconfiguration rather than a kernel-specific recurrence unit [8]. VEXP adds BF16 scalar/SIMD EXP instructions to RISC-V while software FREP, SSR, and

SIMD support handles the remaining partial Softmax [9]. Titopoulos et al. vectorize FlashAttention entirely with standard RVV and software-controlled exponential computation, without custom hardware or scalar operations [10]. These approaches retain programmability at the ISA or processor level, whereas our boundary places scalar recurrence control in a separate SMU and leaves the existing RVV datapath for $O[D]$.

Tightly coupled kernel accelerators target a larger computational unit. RedMule couples a semi-systolic FP16 matrix-multiplication engine to a PULP cluster and shared TCDM, covering complete GEMM kernels [11]. OpenGeMM combines a parameterized GEMM array, lightweight RISC-V control, and a tightly coupled scratchpad with prefetch and buffering mechanisms [12]. Both works optimize complete matrix or GEMM execution and its data movement. Their accelerators therefore differ from a design that isolates state-dependent recurrence work and hands regular $O[D]$ arithmetic back to an existing vector core.

Among the reviewed designs, we did not identify an architecture that dedicates hardware only to the state-dependent online-merge scalar recurrence and, after completion, reuses an existing RVV datapath through a TCDM-mediated weight handoff for the regular $O[D]$ update. To our knowledge, this selective-offload boundary is distinct from prior full-Softmax or full-attention accelerators, programmable-vector or ISA-centered approaches, and tightly coupled full-kernel accelerators.

VI. CONCLUSION

The SMU moves online-softmax merge state updates into a compact TCDM-connected engine while preserving the Spatz ISA and vector pipeline. Measured B3 results outperform the RTL-aligned RVV-assisted baseline at all 16 scaling coordinates and retain component throughput during a shared-TCDM experiment. The measured FSM phases identify vector streaming and command overhead as the next design targets. These results establish cluster-local merge offload as a practical building block for future tiled-attention systems.

REFERENCES

- [1] M. Milakov and N. Gimelshein, "Online normalizer calculation for softmax," *arXiv preprint arXiv:1805.02867*, 2018.
- [2] T. Dao, D. Y. Fu, S. Ermon, A. Rudra, and C. Ré, "Flashattention: Fast and memory-efficient exact attention with io-awareness," in *Advances in Neural Information Processing Systems*, vol. 35, 2022, pp. 16 344–16 359.
- [3] M. Perotti, S. Riedel, M. Cavalcante, and L. Benini, "Spatz: Clustering compact risc-v-based vector units to maximize computing efficiency," *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, vol. 44, no. 7, pp. 2488–2502, 2025.
- [4] J. R. Stevens, R. Venkatesan, S. Dai, B. Khailany, and A. Raghunathan, "Softmax: Hardware/software co-design of an efficient softmax for transformers," in *2021 58th ACM/IEEE Design Automation Conference (DAC)*, 2021, pp. 469–474.
- [5] T. Xia and S. Q. Zhang, "Hyft: A reconfigurable softmax accelerator with hybrid numeric format for both training and inference," in *Proceedings of the 29th ACM/IEEE International Symposium on Low Power Electronics and Design*, 2024, pp. 1–6.
- [6] A. Belano, Y. Tortorella, A. Garofalo, L. Benini, D. Rossi, and F. Conti, "A flexible template for edge generative ai with high-accuracy accelerated softmax and gelu," *IEEE Journal on Emerging and Selected Topics in Circuits and Systems*, vol. 15, no. 2, pp. 200–216, 2025.
- [7] K. Alexandridis, V. Titopoulos, and G. Dimitrakopoulos, "Low-cost flashattention with fused exponential and multiplication hardware operators," in *2025 IEEE Computer Society Annual Symposium on VLSI (ISVLSI)*, 2025, pp. 1–6.
- [8] M. Perotti, M. Raeber, M. Sinigaglia, M. Cavalcante, D. Rossi, and L. Benini, "Spatzformer: An efficient reconfigurable dual-core risc-v v cluster for mixed scalar-vector workloads," in *2024 IEEE 35th International Conference on Application-specific Systems, Architectures and Processors (ASAP)*, 2024, pp. 172–173.
- [9] R. Wang, G. Islamoglu, A. Belano, V. Potocnik, F. Conti, A. Garofalo, and L. Benini, "Vexp: A low-cost risc-v isa extension for accelerated softmax computation in transformers," in *2025 IEEE 32nd Symposium on Computer Arithmetic (ARITH)*, 2025, pp. 37–44.
- [10] V. Titopoulos, K. Alexandridis, and G. Dimitrakopoulos, "Vectorized flashattention with low-cost exponential computation in risc-v vector processors," *The Journal of Supercomputing*, vol. 82, no. 4, 2026.
- [11] Y. Tortorella, L. Bertaccini, D. Rossi, L. Benini, and F. Conti, "Redmule: A compact fp16 matrix-multiplication accelerator for adaptive deep learning on risc-v-based ultra-low-power socs," in *2022 Design, Automation & Test in Europe Conference & Exhibition (DATE)*, 2022, pp. 1099–1102.
- [12] X. Yi, R. Antonio, J. Dumoulin, J. Sun, J. Van Delm, G. P. Paim, and M. Verhelst, "Opengemm: A high-utilization gemm accelerator generator with lightweight risc-v control and tight memory coupling," in *Proceedings of the 30th Asia and South Pacific Design Automation Conference*, 2025, pp. 1055–1061.